

Auto-generated xdcomms system model and specification

Columbia University, as part of the GAPS-CLOSURE project

March 21, 2023

Contents

1	Introduction	2
1.1	About xdcomms and the GAPS-CLOSURE project	2
1.2	About this document	2
2	Summary	2
2.1	High-level overview of xdcomms auto-generation	2
2.2	Goals of xdcomms design and implementation	2
3	Information model	2
3.1	Summary diagram	2
3.2	Primitives and aliases	3
3.3	Cross-domain data format	3
3.4	The hardware interface	5
3.5	The network abstraction	6
3.6	The application model	6
4	Behavioral specification	7
4.1	Control flow	7
4.2	Function-level contracts	7
5	Generator operation	8
5.1	Inputs: The GEDL	8
5.2	Outputs	8
5.3	Whole-system correctness properties and proofs	8
6	References	8

1 Introduction

1.1 About xdcomms and the GAPS-CLOSURE project

==TODO==

1.2 About this document

==TODO==

2 Summary

2.1 High-level overview of xdcomms auto-generation

==TODO==

2.2 Goals of xdcomms design and implementation

==TODO==

3 Information model

In this section we formally model all of the datatypes and classes used in the xdcomms system, and provide intuition as to their purpose; the application control flow and function behavior is fully specified in Section 4.

3.1 Summary diagram

The following UML diagram provides a bird's eye summary of the information model.
==TODO==

3.2 Primitives and aliases

The following aliases are used in our information model.

==TODO== (Section 3.3).

```
MuxTag :: int
SecTag :: int
TypTag :: int
GTag :: (MuxTag, SecTag, TypTag)
```

An Enclave is a name for a set of devices operating in a closed network at the same security level. A Level is a name given to a security level for an application endpoint, which is used to specify what data the endpoint can receive from which other endpoints, and across which communication channels.

```
Enclave :: string
Level :: string
```

==TODO==

```
DataType :: string
Status :: int
RequestID :: int
ResponseID :: int
Timestamp :: int
```

==TODO== (Section 3.4).

```
AppEndpt :: URI
DevType :: string
FHdl :: int
TagDev :: (GTag, Device)
```

3.3 Cross-domain data format

This subsection describes the format of program data which can be transferred from one application endpoint to another via xdcomms, and the intermediate representations of said data.

3.3.1 The Data type

xdcomms uses a simple type, `Data`, to capture a collection of program data which is eligible to be transferred cross-domain to another application endpoint. It is isomorphic to a struct in C; the purpose is to *lift* valid, serializable structs into the xdcomms informaton model.

```
type Data =
  dtype      :: DataType  // string identifier for this data format.
  contents   :: [CValue]  // ordered list of fields, where CValue refers to
                           // a typed value in C that is a scalar, vector,
                           // or serializable struct.
```

We restrict `CValue` to *serializable* datatypes, meaning a `CValue` is not a pointer, or, if it is a struct, the struct fields are all serializable (i.e. not pointers). Because application endpoints are physically and memory isolated, pointers cannot be transferred cross-domain. Fixed-length arrays, however, are valid `CValues` (because fixed-length arrays are stored by value in C structs).

3.3.2 The Marshalled type

Serialization of data into a packet requires awareness of the network and the protocol between the communicating devices. There is, however, device-independent information that must be appended to the data being transferred; we refer to this process as *marshalling*.

The `Marshalled` type captures data which is ready for serialization; it contains the original data, along with a *trailer* that stores a sequence number and error-correction codes:

```
type Marshalled =
  mcontents :: Data
  trailer   :: Trailer

type Trailer =
  seq :: uint32_t
  rqr :: uint32_t
  oid :: uint32_t
  mid :: uint16_t
  crc :: uint16_t
```

`Marshalled` and `Trailer` are *packed* so that their fields have no padding, for ease of serialization. ==PERATON: What are the other trailer fields for?==

3.3.3 The Serialized type

Before being sent cross-domain, data converted into an in-memory **Marshaled** struct with a sequence number and error correction codes must additionally be:

- Tagged with a **GTag**.
- Given a unique ID to distinguish it from other incoming request or response packets at the destination.
- Serialized into a stream of bytes in network order.
- Coerced into the packet format expected at the destination device.

This transformation is mediated through the **Serialized** type:

```
type Serialized =  
    tag    :: GTag  
    bytes  :: [uint8_t]  
    reqid  :: RequestID  
    rspid  :: ResponseID
```

The GAPS tag, of type **GTag**, is a tuple of three tags: a **MuxTag** identifying the source and destination applications, a **SecTag** identifying the security levels between which the data is being transferred, and a **TypTag** encoding the **DataType** as an integer (so that it can be unmarshalled at the destination). So for **GTag** we have the type **GTag** : (**MuxTag**, **SecTag**, **TypTag**).

The request ID **reqid** is set if this communication is a request, and the response ID **rspid** is set if this communication is a reply to a request. ==PERATON: Where are the request and response ID set? How exactly are they used?==

Conversion of a **Marshaled** struct proceeds in several steps. First, the **Marshaled** field is copied into a parallel (packed) struct, in which every field (including those in the trailer) is represented as bytes in *network byte order*. This requires that floats be converted into a device-independent bytestring format. Then, this parallel struct is cast to a void pointer and treated as an array of bytes, stored in the bytes field of the **Serialized** struct. ==PERATON: which part of this is device-aware? The codec functions don't seem to perform any device-specific packetizing. Is this hiding in zmq send?==

Conversion between a **Marshaled** struct and its corresponding **Serialized** struct is mediated by the **serialize()** and **deserialize()** functions in the device-aware **Codec** class (Section 3.4).

3.4 The hardware interface

==TODO==

3.4.1 The HALConfig class

==TODO==

3.4.2 The Device class

==TODO==

3.4.3 The Codec class

==TODO==

3.5 The network abstraction

==TODO==

3.5.1 The Binding type

==TODO==

3.5.2 The XDContext type

==TODO==

3.5.3 The Wrapper class

==TODO==

3.5.4 The Handler class

==TODO==

3.6 The application model

==TODO==

3.6.1 The MasterSequence class

==TODO==

3.6.2 The EventQueue class

==TODO==

3.6.3 The RPCTransaction class

==TODO==

3.6.4 The HAL class

==TODO==

3.6.5 The App class

==TODO==

3.6.6 The AppThread class

==TODO==

4 Behavioral specification

==TODO==

4.1 Control flow

==TODO (diagram)==

4.2 Function-level contracts

==TODO==

4.2.1 `Wrapper.marshall()`

The `marshall()` function is called by the `Wrapper` to marshall the incoming arguments (the arguments to the cross-domain function in the original, unpartitioned program) into one serializable data structure with error correction and a sequence number.

Arguments: `d :: Data`, `req_counter: int` (`d` is unpacked and may represent multiple or no arguments)

Return value: `m :: Marshalled`

Pre-conditions:

- The value of `req_counter` is either non-negative or `INT_MIN`.
- The `Marshalled` struct defines, in the same order as the arguments to `marshall()`, one field of the same type and name as each argument, followed by a `Trailer`.

Post-conditions:

- For each incoming argument, the corresponding field in `m` is set to the value of that argument.
- The trailer of `m` has its `seq` field set to the value of `req_counter`.

5 Generator operation

==TODO==

5.1 Inputs: The GEDL

==TODO==

5.2 Outputs

5.3 Whole-system correctness properties and proofs

==TODO==

6 References